



AI Assistance in Medical Imaging using PyTorch

Abhishek Kumar

Deep Learning Engineer @Predible Health

Aditya Bagari


Undergraduate @IITMadras

Deep learning is everywhere!

- Presence in:
 - Image Recognition
 - Voice Recognition
 - Machine Translation
 - Self-driving cars etc..
- Artificial Neural Networks are the engines
- Inspired by biological neural networks




What is a neural network?



Biological Neural Networks


- Group of connected biological nerve cells
- Human brain is a biological neural network



Artificial Neural Networks



- Group of connected artificial neurons
- Mathematical approximation of biological NNs

What is an artificial neuron?



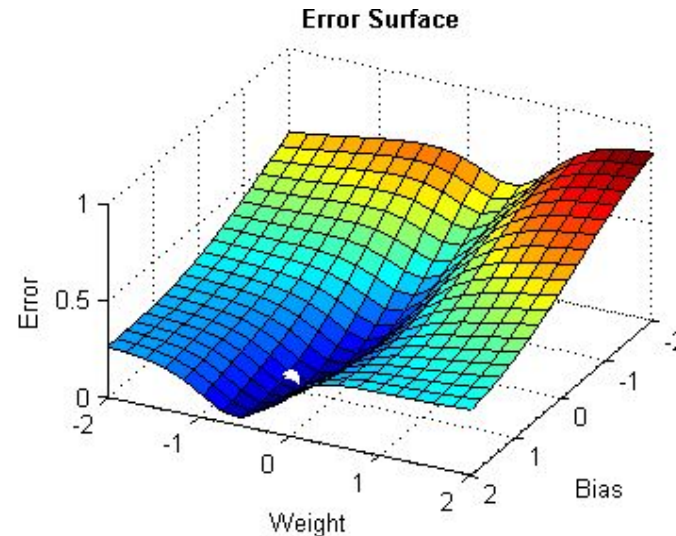

- Parameters to be learnt – weights and biases

Multilayer Perceptron



How do the networks learn?

- Cost function – compute error between prediction and label
- Learning by gradient descent – adjust parameters to reduce cost
- **Backpropagation** – compute gradients by applying chain rule of differentiation
- Gradient descent - Blind man walking down a hill step-by-step in search of a valley



Firing up your Jupyter Notebooks!

- `conda create env -f med-torch.yml`
- `conda activate med-torch / source activate med-torch`


PYTORCH

Introduction to PyTorch – hands on

- Tensor handling ops – matrix multiplication, addition etc..
- Define layers
- Neural network modules
- Forward pass
- Loss computation
- Optimization step
- NumPy to Torch tensor conversion


Applying deep learning on hand-written digits

Images are numbers too!



Importance of normalization:

- “Centers” the data
- Stabilizes weight updates




MNIST demo – hands on

- Dataset of 60,000 images
- 10 labels [0-9]
- 28x28 images
- Multi-layer Perceptron
- Learning by gradient descent



Classify digits using a neural network

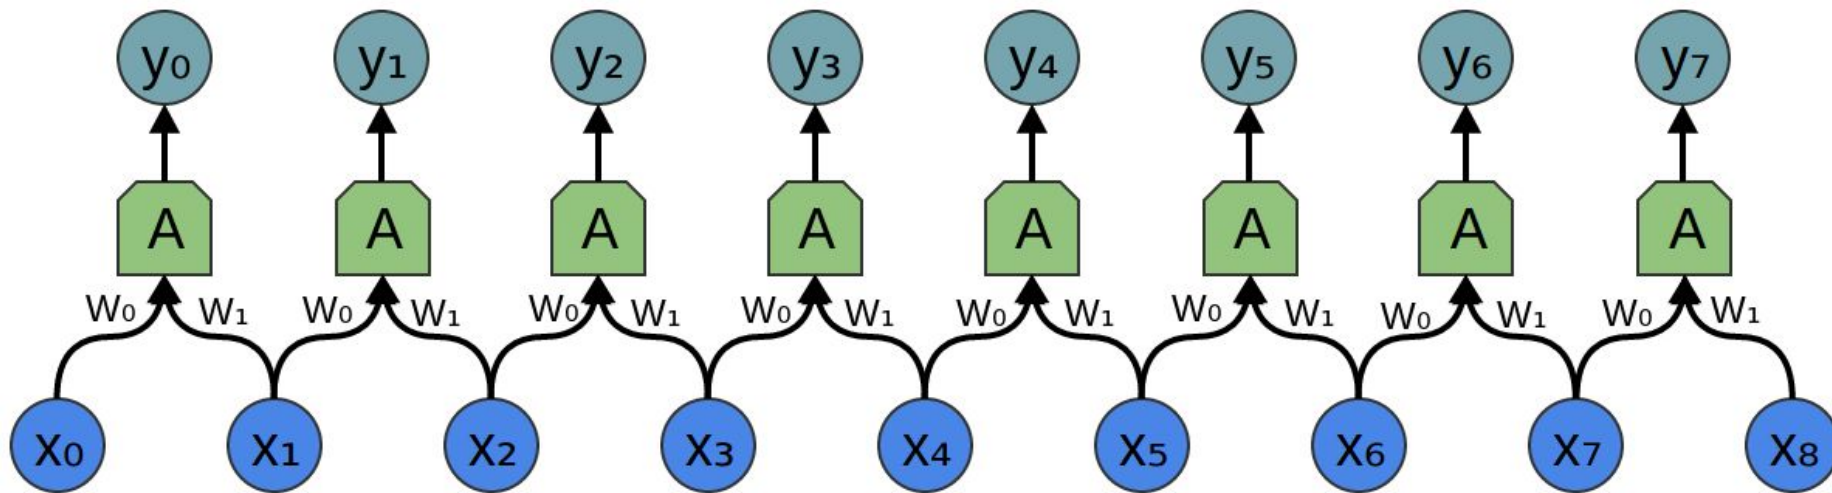
The problem of overfitting



- Too many parameters can result in over-fitting
- Too less can result in under-fitting
- Choose the right number of parameters


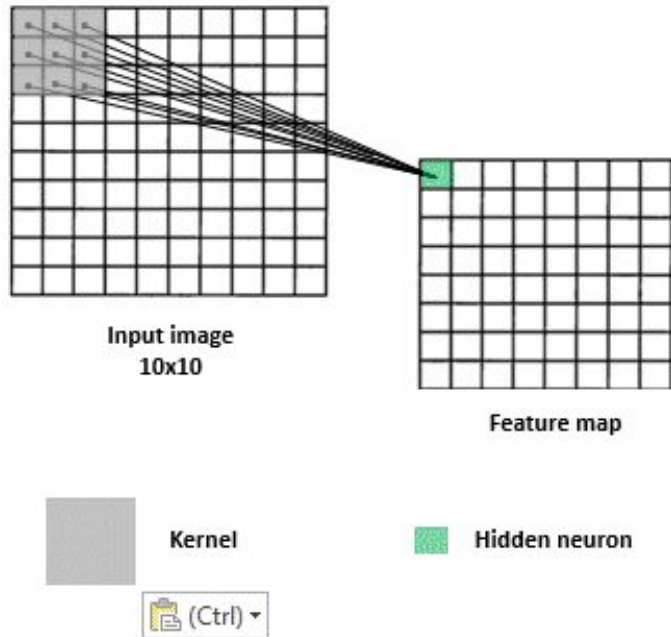
Convolutional Neural Networks

- Weights are shared between neurons – w_0 and w_1
- Prevents networks from overfitting




Understanding convolutions in 2D

- “Roll” the layer over the entire image



The max-pooling layer

- Take only the maximum defined by the kernel size
- Helps in reducing the size of the network
- Retains important features



Build a Convolutional Neural Network

- Put together a few convolutions, max-poolings and fully connected layers.


```
class MnistModel(nn.Module):
    def __init__(self):
        super(MnistModel, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 5, padding=2)
        self.conv2 = nn.Conv2d(32, 64, 5, padding=2)
        self.fc1 = nn.Linear(64*7*7, 1024)
        self.fc2 = nn.Linear(1024, 10)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), 2)
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, 64*7*7) # reshape Variable
        x = F.relu(self.fc1(x))
        x = F.dropout(x, training=self.training)
        x = self.fc2(x)
        return F.log_softmax(x)
```

MNIST demo with CNNs – hands on

Pixel-wise segmentation

- Classify every pixel in the image by “rolling” the network on patches
- Process 32x32 patches with stride 1



WiFi : pycon_workshop
Password: wspycon2018